

CSE
2021

Computer
Organization

CPU DESIGN

Pipelining

Prof. H. Roumani,
Dept of CS & Eng, York University

PIPELINING

- Basic Concepts
- Implementation
- Hazards
- Hazard Handling

The Idea

Assume delays of 1, 2, and 4 min for ordering, paying, and making a sandwich in the worst case.

What is the latency and throughput of a “single-cycle” restaurant (door opens and shuts at regular intervals allowing one customer to enter and one to leave)?

7 min, 8.6 person/hr

Is this arrangement simple? Does it cater to the slowest customer? Is it wasteful (idle employees)?

Yes. Yes. Yes.

The Idea, cont.

What is the latency and throughput of a “pipelined” restaurant?

12 min, 15.0 person/hr

Compared to the single-cycle restaurant, latency has become worse but throughput is better!

Is this arrangement as simple? Caters to the slowest customer? Is it wasteful (idle employees)?

Almost (need to marshal). No. No.

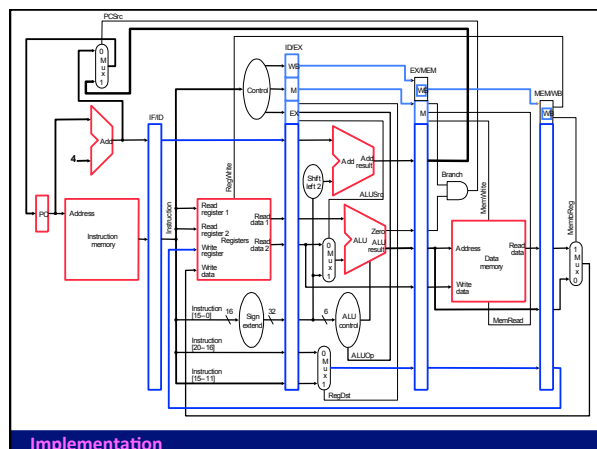
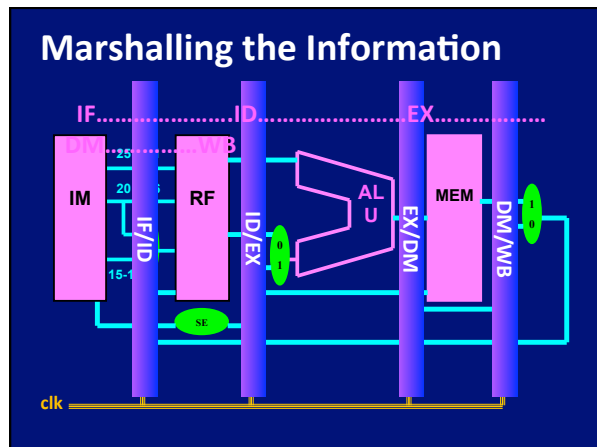

Basic Concepts

- **Assembly Line**
Stages (Order, Pay, Make) and Marshaling.
- **Balancing, Latency, and Throughput**
Balancing the stages increases the latency
- **The Performance Potential**
Up to n-fold improvement in throughput (for n stages)
- **No Dependencies**
 - Can I add small fries?
 - I'll have the same as my friend in the front.

A Pipelined CPU

- **Assembly Line**
Five stages:
IF (in **IM**), ID (in **RF**), EX (in **ALU**), DM (in **MEM**), WB (in **RF**)
- **Balancing, Latency, and Throughput**
IF (200), ID (50), EX (100), DM (200), WB (50)
Balance 200. Latency is thus 1000 (not 600) or 1 GHz
- **The Performance Potential**
Up to 5 GHz

Pipelines



Exercise

During a particular clock cycle, the five highlighted instructions were be in the pipeline; in particular, the addi instruction of address 420 was in its EXecution stage. Determine the content of EX/MEM at the end of that cycle.

```
400 addi $a0, $0, 25
404 addi $t6, $0, 13
408 addi $s0, $0, 47
412 add $t5, $0, $0
416 sub $t0, $0, $a0
420 addi $t6, $a0, 22
424 lw $t7, 24($a0)
428 lw $t8, 12($a0)
...
```

FIELD	WIDT H	VALUE
WB		
M		
Add Result		
Zero		
ALU Result		
Read Data 2		
Write Register		

Answer

During a particular clock cycle, the five highlighted instructions were be in the pipeline; in particular, the addi instruction of address 420 was in its EXecution stage. Determine the content of EX/MEM at the end of that cycle.

```
400 addi $a0, $0, 25
404 addi $t6, $0, 13
408 addi $s0, $0, 47
412 add $t5, $0, $0
416 sub $t0, $0, $a0
420 addi $t6, $a0, 22
424 lw $t7, 24($a0)
428 lw $t8, 12($a0)
...
```

FIELD	WIDT H	VALUE
WB	2	10 ₈
M	3	000 ₈
Add Result	32	512 ₀
Zero	1	0
ALU Result	32	47 ₀
Read Data 2	32	13 ₀
Write Register	5	14 ₀

Exercise

During a particular clock cycle, the five highlighted instructions were be in the pipeline; in particular, the addi instruction of address 420 was in its EXecution stage. Determine the values of the wires at the end of that cycle (see next slide).

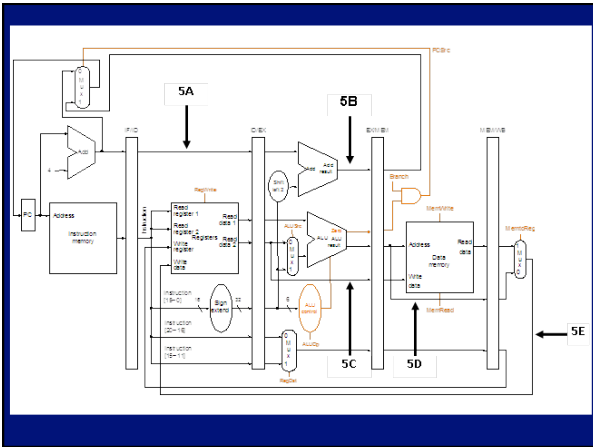
```
400 addi $a0, $0, 25
404 addi $t6, $0, 13
408 addi $s0, $0, 47
412 add $t5, $0, $0
416 sub $t0, $0, $a0
420 addi $t6, $a0, 22
424 lw $t7, 24($a0)
428 lw $t8, 12($a0)
...
```

FIELD	VALUE
5A	
5B	
5C	
5D	
5E	

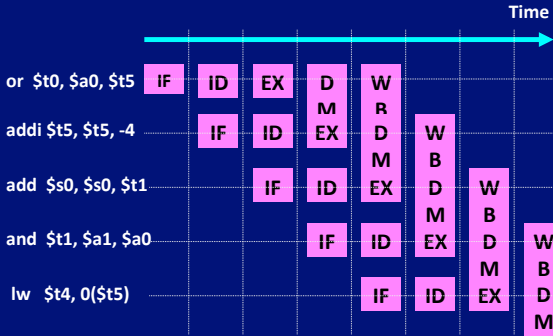
Answer


During a particular clock cycle, the five highlighted instructions were be in the pipeline; in particular, the addi instruction of address 420 was in its EXecution stage. Determine the values of the wires at the end of that cycle (see next slide).

400 addi \$a0, \$0, 25	FIELD	VALUE
404 addi \$t6, \$0, 13	5A	428
408 addi \$s0, \$0, 47	5B	512
412 add \$t5, \$0, \$0	5C	13
416 sub \$t0, \$0, \$a0	5D	-25
420 addi \$t6, \$a0, 22	5E	0
424 lw \$t7, 24(\$a0)		
428 lw \$t8, 12(\$a0)		
...		



Perfect world: no dependencies





Hazards

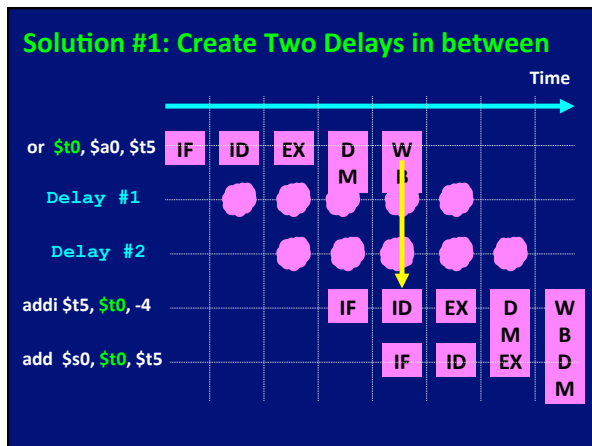
Types, Detection, and Avoidance

- **Structural**
Can I add small fries (out of sequence)
Out of sequence like addm. Never happens in MIPS
- **Data**
I order based on how much my friend paid
RF/ALU and ALU/DM
- **Control**
I'll have the same as my friend in the front
All branches (conditional transfers)

Type 1 Data Hazard: Compute + dependency

Example:

```
addi    $t0, $0, 15
or      $s5, $s5, $t0
```



But how to populate the delay slots?

1. Nop's

Have the compiler insert two nop's

BAD! Creates DRAM access delays and crowds the cache

2. Benign instructions

Have the compiler reorder the code

May or may not succeed

3. Bubbles

Have the CPU stall the pipeline

OK but we lost ~half the gain!

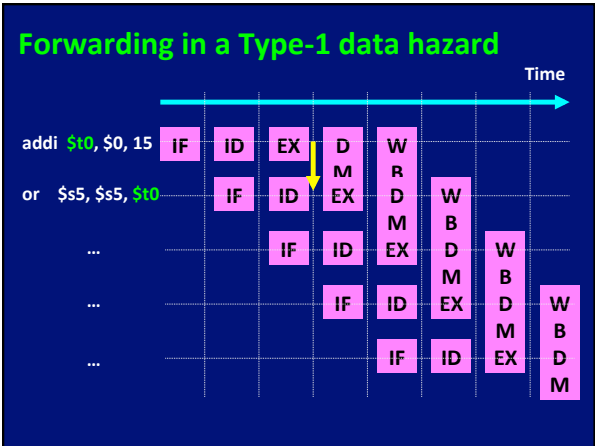
None of these is good ☹


Solution #2: FORWARDING

Change the datapath so that the output of the ALU is sent back to its input, i.e. forwarded directly to the next instruction.

Data is thus available just in time rather than having to wait for DM and WB.

This is a perfect solution! No stalls!





Type 2 Data Hazard:
Load + dependency

Example:

```
lw    $t0, 16($0)
or     $s5, $s5, $t0
```

Type 2 Data Hazard


Solution #1

We can, as in Type-1, insert two delays between the dependent instructions. But this will degrade the performance.

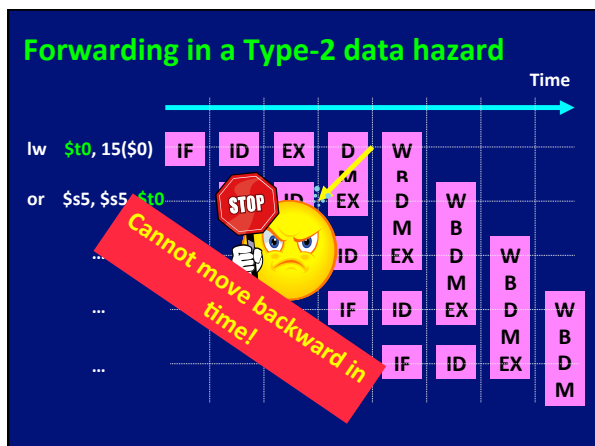
Solution #2

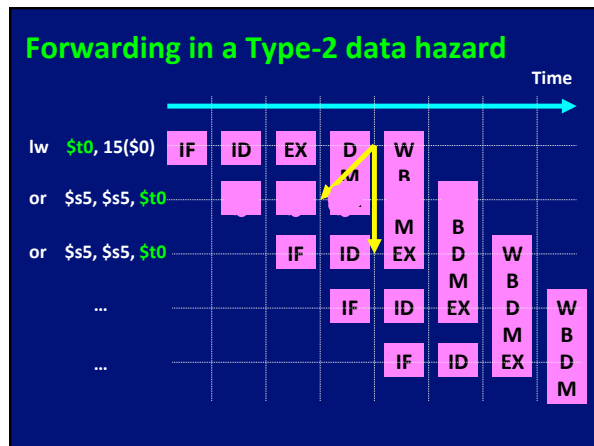
Forwarding worked very well for us in Type-1. Can we employ it here?

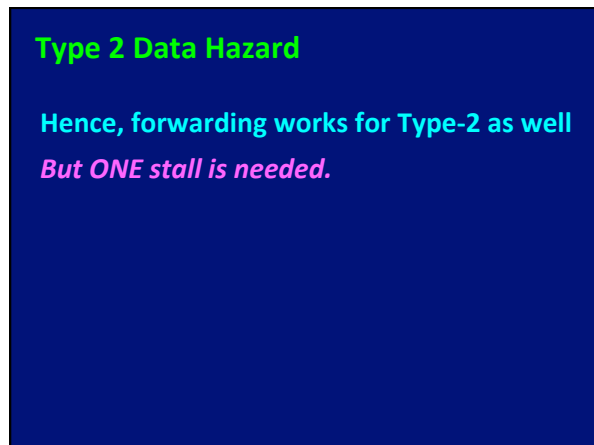
Forwarding in a Type-2 data hazard

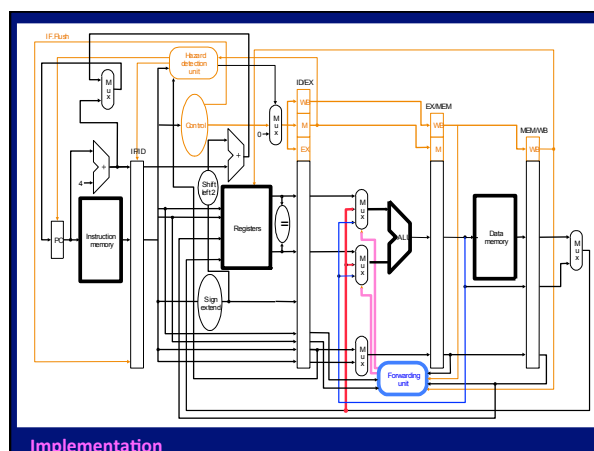


Forwarding in a Type-2 data hazard





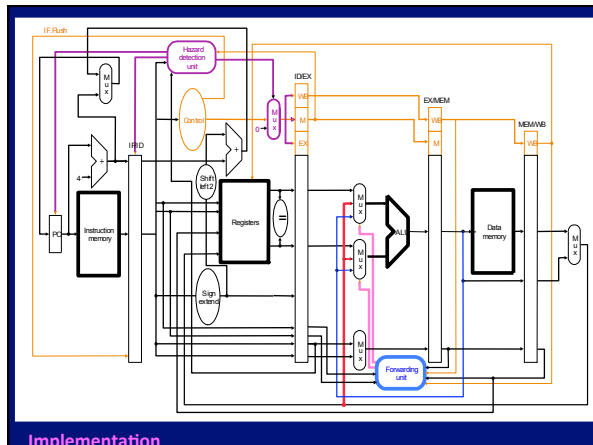




But how can we create a stall?

When a type-2 is detected (i.e. lw in EX and some ins that depends on it in ID), create a stall between them:

- The instruction in ID must stay in ID. How?
Disable IF/ID!
- The instruction in IF must stay in IF. How?
Disable PC!
- A fake but benign instruction must be marshaled from ID to EX. How?
Zero out the controls in ID/EX



Implementation

Control Hazards: branches

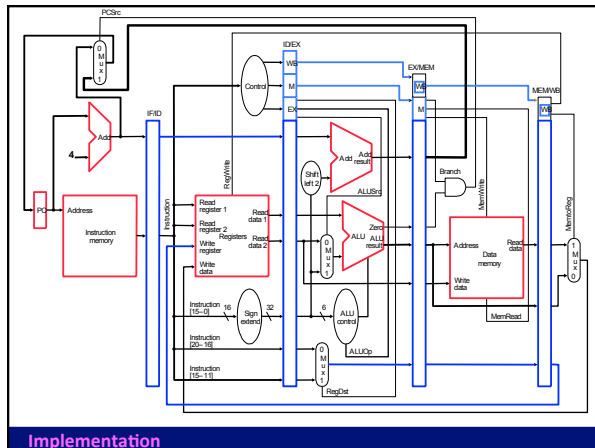
By the time the branch decision is taken, two instructions are already in the pipeline! *Why?*

Example:

```
beq    $t0, $s0, loop
or     $s5, $s5, $a0
addi   $t3, $t3, 1
```

Hence, two stalls are needed per branch ☹

Pipelines

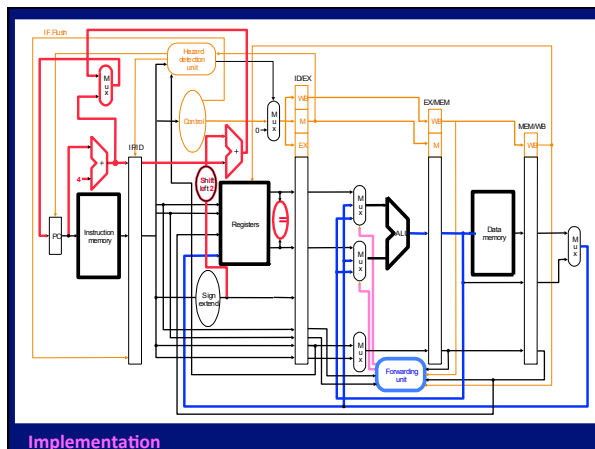


Control Hazards ...

Branches account for ~25% of all instructions
 ➔ two stalls per branch is unacceptable.

Reduce them to one stall by detecting equality and computing the target in ID rather than EX

To implement this stall, flush the IF/ID register. When the clock ticks, IF/ID would hold nop, and IF would be fetching PC+4 or PC+4+target*4 (PC=branch address).



Approaches to Branch Hazards

- Always stall once
- Assume the branch will not be taken
- Predict the branch target
- Delay the branch (S/W implication)

Critique these approaches.

Assume a 100-iteration loop and analyze the factors that determine the likely penalty per branch in each approach

The Pipeline

Performance

```
do
{
    int e1 = ar[i];
    sum += e1;
    sha ^= e1;
    i++;
} while (e1 != last);
```

```
loop: lw    $t5, ar($t0)
      addu  $s0, $s0, $t5
      xor   $s2, $s2, $t5
      addi  $t0, $t0, 4
      bne   $t5, $a0, loop
```

How long does it take to execute this loop on a **single-cycle** machine?

Assume component latencies of:
RF=50, ALU=100, and MEM (both IM and DM)=200 ps

Answer: 50 ns/iteration

```
do
{
    int e1 = ar[i];
    sum += e1;
    sha ^= e1;
    i++;
} while (e1 != last);
```

```
loop: lw    $t5, ar($t0)
      addu  $s0, $s0, $t5
      xor   $s2, $s2, $t5
      addi  $t0, $t0, 4
      bne   $t5, $a0, loop
```

How long does it take to execute this loop on a multi-cycle machine?

Assume component latencies of:
RF=50, ALU=100, and MEM (both IM and DM)=200 ps

Answer: 2.5 ns/iteration

```
do
{
    int e1 = ar[i];
    sum += e1;
    sha ^= e1;
    i++;
} while (e1 != last);
```

```
loop: lw    $t5, ar($t0)
      addu  $s0, $s0, $t5
      xor   $s2, $s2, $t5
      addi  $t0, $t0, 4
      bne   $t5, $a0, loop
```

How long does it take to execute this loop on a pipelined machine?

Assume component latencies of:
RF=50, ALU=100, and MEM (both IM and DM)=200 ps

Answer: 1.0 to 1.6 ns/iteration